

Download a the dataset sample from Supervisely

```
#@title Download a the dataset sample from Supervisely
```

```
!curl https://assets.supervisely.com/supervisely-supervisely-assets-public/teams_storage/W/6/pt/ANAGVgKaC62tTrDQWK5JhNP2dd8ynqaTKSM1QdVoA
```

```

→ % Total    % Received % Xferd  Average Speed   Time    Time       Time  Current
   100    548M    100    548M    0      0  8892k      0  0:01:03  0:01:03  --:--:--  12.7M

```

Extract the tarball 'mvtec.tar' and look for test images for the hazelnut subset

```
#@title Extract the tarball 'mvtec.tar' and look for test images for the hazelnut subset
```

```
!tar -xf mvtec.tar
```

```
!ls test/img/hazelnut_*
```

```

→ tar: Removing leading `/' from member names
test/img/hazelnut_crack_002.png test/img/hazelnut_print_004.png
test/img/hazelnut_cut_003.png test/img/hazelnut_print_005.png
test/img/hazelnut_good_023.png test/img/hazelnut_print_006.png
test/img/hazelnut_good_037.png test/img/hazelnut_print_007.png
test/img/hazelnut_hole_005.png test/img/hazelnut_print_009.png
test/img/hazelnut_hole_010.png test/img/hazelnut_print_011.png
test/img/hazelnut_hole_013.png test/img/hazelnut_print_012.png
test/img/hazelnut_hole_016.png test/img/hazelnut_print_013.png
test/img/hazelnut_print_003.png test/img/hazelnut_print_016.png

```

```
!git clone https://github.com/akridata-ai/ZS-CLIP-AC-naive.git
```

```

→ Cloning into 'ZS-CLIP-AC-naive'...
remote: Enumerating objects: 16, done.
remote: Counting objects: 100% (16/16), done.
remote: Compressing objects: 100% (14/14), done.
remote: Total 16 (delta 2), reused 12 (delta 0), pack-reused 0 (from 0)
Receiving objects: 100% (16/16), done.
Resolving deltas: 100% (2/2), done.

```

```
%cd ZS-CLIP-AC-naive
```

```
→ /content/ZS-CLIP-AC-naive/ZS-CLIP-AC-naive
```

```
!git branch -a
```

```

→ * main
   remotes/origin/HEAD -> origin/main
   remotes/origin/feature/template-code
   remotes/origin/main

```

```
!git checkout -b feature/template-code origin/feature/template-code
```

```

→ Branch 'feature/template-code' set up to track remote branch 'feature/template-code' from 'origin'.
Switched to a new branch 'feature/template-code'

```

```
!git pull
```

```
→ Already up to date.
```

```
!git config --global user.email "amalamahalakshmi@email.com"
```

```
!git config --global user.name "amalamahalakshmi"
```

```
%cd /content/ZS-CLIP-AC-naive
```

```
→ /content/ZS-CLIP-AC-naive
```

```
!ls
```

```

→ clip_ac.py  meta.json  README.md      test
  data       mvtec.tar  requirements.txt train
  LICENSE.md __pycache__ spec.py        ZS-CLIP-AC-naive

```

```
!cp /content/test/img/hazelnut_* data/
```

```
!pip install -qr requirements.txt
```

```
→ Preparing metadata (setup.py) ... done
```

```

%%writefile spec.py
"""
Spec for zero-shot defect classification with CLIP.
"""
from pydantic import BaseModel
from typing import List, Dict

class DefectClassificationSpec(BaseModel):
    # final labels you want in the confusion matrix (order is fixed)
    class_names: List[str] = ["good", "print", "hole", "cut", "crack"]

    # base CLIP model (try ViT-B/16 first; switch to ViT-L/14 if GPU has room)
    model_name: str = "ViT-B/16"

    # two-stage: first 'good' vs 'defect', then if defect -> {print, hole, cut, crack}
    use_two_stage: bool = True

    # simple test-time augmentation (original + horizontal flip)
    use_tta: bool = True

    # temperature (logit scaling); 50-100 works well for CLIP
    temperature: float = 100.0

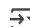
    # if two-stage, confidence margin to call something 'good'
    # larger = stricter 'good' (reduces false goods)
    good_margin: float = 0.0

    # natural-language phrases per class (richer semantics than single words)
    phrases: Dict[str, List[str]] = {
        "good": [
            "a clean, defect-free hazelnut",
            "an undamaged hazelnut",
            "a perfect normal hazelnut without any defects",
            "an intact hazelnut with smooth surface",
            "a pristine hazelnut"
        ],
        "print": [
            "a hazelnut with a print defect",
            "a hazelnut with printed text on the shell",
            "a hazelnut with a stamped marking",
            "a hazelnut with surface printing marks",
            "a hazelnut with an ink print on the shell"
        ],
        "hole": [
            "a hazelnut with a hole defect",
            "a hazelnut with a drilled hole",
            "a hazelnut with a punctured hole in the shell",
            "a hazelnut with a perforated shell",
            "a hazelnut with a hole on the surface"
        ],
        "cut": [
            "a hazelnut with a cut defect",
            "a sliced hazelnut",
            "a hazelnut with a cut surface",
            "a hazelnut cut in half",
            "a hazelnut showing a clean cut"
        ],
        "crack": [
            "a hazelnut with a crack defect",
            "a cracked hazelnut shell",
            "a broken hazelnut with a visible crack",
            "a hazelnut shell split open",
            "a hazelnut showing a fracture"
        ]
    }

    # extra templates (from CLIP/ImageNet style) - multiplied with phrases
    templates: List[str] = [
        "a photo of {}",
        "a close-up photo of {}",
        "a cropped photo of {}",
        "a bright photo of {}",
        "a low-light photo of {}",
        "a high-resolution photo of {}",
        "a clean product photo of {}",
        "an image of {}",
        "a centered photo of {}",
        "a detailed photo of {}",
        "a studio photo of {}",
        "a DSLR photo of {}"
    ]

```

```
# prompts for stage-1 binary decision
stage1_good_phrases: List[str] = [
    "a clean, defect-free hazelnut",
    "a normal undamaged hazelnut",
    "a pristine hazelnut with no defects",
]
stage1_defect_phrases: List[str] = [
    "a defective or damaged hazelnut",
    "a hazelnut with some kind of defect",
    "a faulty hazelnut with visible damage",
]
```

 Overwriting spec.py

```
%%writefile clip_ac.py
"""
Zero-shot defect classification using CLIP with:
- prompt ensembling over templates & phrases
- optional two-stage head (good vs defect, then defect type)
- light TTA (horizontal flip)
"""

from pathlib import Path
from typing import List, Tuple, Dict

import clip
import torch
from PIL import Image, ImageOps

from spec import DefectClassificationSpec

def _build_prompts(phrases: List[str], templates: List[str]) -> List[str]:
    # e.g., template="a photo of {}", phrase="a hazelnut with a crack defect"
    return [tpl.format(ph) for tpl in templates for ph in phrases]

@torch.no_grad()
def _encode_prompt_set(model, device, prompts: List[str]) -> torch.Tensor:
    # tokenize, encode, L2-normalize, then average across prompts
    tokens = clip.tokenize(prompts).to(device)
    feats = model.encode_text(tokens)
    feats = feats / feats.norm(dim=-1, keepdim=True)
    mean = feats.mean(dim=0, keepdim=True)
    mean = mean / mean.norm(dim=-1, keepdim=True)
    return mean # [1, d]

@torch.no_grad()
def _encode_classes(model, device, class_names: List[str],
                    phrases_map: Dict[str, List[str]],
                    templates: List[str]) -> torch.Tensor:
    # returns [C, d] class text embeddings
    cls_feats = []
    for cls in class_names:
        phrases = phrases_map[cls]
        prompts = _build_prompts(phrases, templates)
        mean = _encode_prompt_set(model, device, prompts)
        cls_feats.append(mean)
    return torch.cat(cls_feats, dim=0)

def _parse_label(fname: str) -> str:
    # hazelnut_good_023.png -> good; hazelnut_crack_002.png -> crack
    name = fname.lower()
    if "good" in name:
        return "good"
    parts = name.split("_")
    if len(parts) >= 2:
        return parts[1].replace(".png", "")
    return "good"

@torch.no_grad()
def classify_defects(spec: DefectClassificationSpec, test_dir: Path) -> Tuple[List[str], List[str]]:
    device = "cuda" if torch.cuda.is_available() else "cpu"

    # load model with gentle fallback
    try:
        model, preprocess = clip.load(spec.model_name, device=device)
    except RuntimeError:
        try:
            model, preprocess = clip.load("ViT-B/16", device=device)
        except RuntimeError:
            model, preprocess = clip.load("ViT-B/32", device=device)
    model.eval()
```

```

# ----- text features -----
# stage 1: good vs defect (binary)
if spec.use_two_stage:
    stage1_names = ["good", "defect"]
    stage1_phrases = {
        "good": spec.stage1_good_phrases,
        "defect": spec.stage1_defect_phrases,
    }
    stage1_text = _encode_classes(model, device, stage1_names, stage1_phrases, spec.templates) # [2, d]

# stage 2: per-defect classes
text_features = _encode_classes(model, device, spec.class_names, spec.phrases, spec.templates) # [C, d]

y_true, y_pred = [], []

for img_path in sorted(test_dir.glob("*.png")):
    true_label = _parse_label(img_path.name)
    y_true.append(true_label)

# ----- image features with TTA -----
img = Image.open(img_path).convert("RGB")
images = [img]
if spec.use_tta:
    images.append(ImageOps.mirror(img))

img_feats = []
for im in images:
    image = preprocess(im).unsqueeze(0).to(device)
    feat = model.encode_image(image)
    feat = feat / feat.norm(dim=-1, keepdim=True)
    img_feats.append(feat)
image_feat = torch.stack(img_feats, dim=0).mean(dim=0) # [1, d]
image_feat = image_feat / image_feat.norm(dim=-1, keepdim=True)

# ----- two-stage decision (optional) -----
if spec.use_two_stage:
    logits_bin = spec.temperature * (image_feat @ stage1_text.T) # [1, 2]
    probs_bin = logits_bin.softmax(dim=-1).squeeze(0)
    prob_good = float(probs_bin[0].item())
    prob_defect = float(probs_bin[1].item())

    # margin gate: call good only if clearly more probable than defect
    if (prob_good - prob_defect) >= spec.good_margin:
        y_pred.append("good")
        continue # no need stage-2 if good

# ----- stage-2: multi-class among spec.class_names -----
logits = spec.temperature * (image_feat @ text_features.T) # [1, C]
probs = logits.softmax(dim=-1).squeeze(0)
pred_idx = int(probs.argmax().item())
y_pred.append(spec.class_names[pred_idx])

return y_true, y_pred

```

↗ Overwriting clip_ac.py

```

from pathlib import Path
from spec import DefectClassificationSpec
from clip_ac import classify_defects

# 3.1 run classification
spec = DefectClassificationSpec(
    # try a stronger backbone if you have GPU memory:
    # model_name="ViT-L/14",
    use_two_stage=True,
    use_tta=True,
    temperature=100.0,
    good_margin=0.0, # you can try 0.05-0.15 to be stricter for 'good'
)
y_true, y_pred = classify_defects(spec, Path("data"))

# 3.2 evaluation
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
import matplotlib.pyplot as plt

labels = spec.class_names
cm = confusion_matrix(y_true, y_pred, labels=labels)
acc = (np.trace(cm)/cm.sum()) if cm.sum() > 0 else 0.0
print("Labels order:", labels)
print(f"Accuracy: {acc*100:.2f}%\n")
print(classification_report(y_true, y_pred, labels=labels, zero_division=0))


```

```
# 3.3 plot confusion matrix
fig = plt.figure(figsize=(6,5))
plt.imshow(cm, interpolation='nearest')
plt.title("Hazelnut Defect Classification – Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.xticks(range(len(labels)), labels, rotation=45, ha="right")
plt.yticks(range(len(labels)), labels)
for i in range(len(labels)):
    for j in range(len(labels)):
        plt.text(j, i, cm[i, j], ha="center", va="center")
plt.tight_layout()
plt.show()

# 3.4 peek a few predictions
for f, t, p in list(zip(sorted(Path("data").glob("*.png")), y_true, y_pred))[:10]:
    print(f"{f.name:30} true={t:5} pred={p:5}")
```

Labels order: ['good', 'print', 'hole', 'cut', 'crack']
Accuracy: 61.11%

	precision	recall	f1-score	support
good	0.00	0.00	0.00	2
print	0.82	0.90	0.86	10
hole	1.00	0.25	0.40	4
cut	0.00	0.00	0.00	1
crack	0.25	1.00	0.40	1
accuracy			0.61	18
macro avg	0.41	0.43	0.33	18
weighted avg	0.69	0.61	0.59	18



hazelnut_crack_002.png	true=crack	pred=crack
hazelnut_cut_003.png	true=cut	pred=print
hazelnut_good_023.png	true=good	pred=cut
hazelnut_good_037.png	true=good	pred=print
hazelnut_hole_005.png	true=hole	pred=hole
hazelnut_hole_010.png	true=hole	pred=crack
hazelnut_hole_013.png	true=hole	pred=crack
hazelnut_hole_016.png	true=hole	pred=crack
hazelnut_print_003.png	true=print	pred=print
hazelnut_print_004.png	true=print	pred=print

```
import matplotlib.pyplot as plt
from PIL import Image
from pathlib import Path

# filter only correct predictions
correct = [(f, t, p) for f, t, p in zip(sorted(Path("data").glob("*.png")), y_true, y_pred) if t == p]

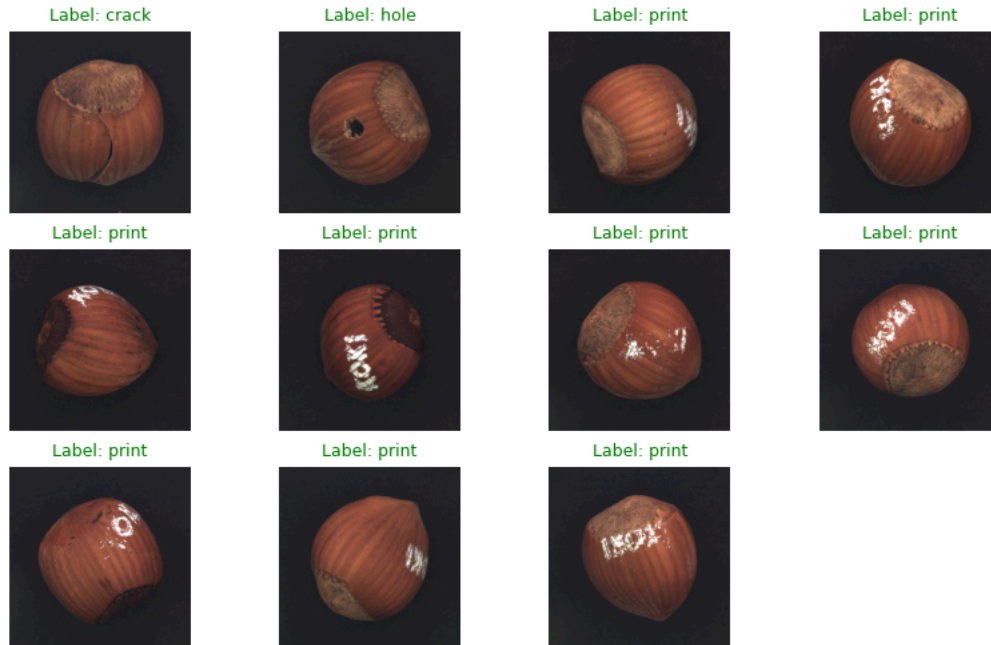
def show_correct(samples):
    plt.figure(figsize=(10, 6))
    plt.suptitle("Predicted Defects in Images", fontsize=16, fontweight="bold")
    for i, (f, t, p) in enumerate(samples[:12]): # show max 12
        img = Image.open(f)
```

```
plt.subplot(3, 4, i+1)
plt.imshow(img)
plt.axis("off")
plt.title(f"Label: {t}", fontsize=9, color="green")
plt.show()
```

```
# Show only correct predictions
show_correct(correct)
```



Predicted Defects in Images



Start coding or [generate](#) with AI.